

Pipelining Tools for HPC Workflows

Using Bash, Snakemake and Nextflow

Sam Friedman

Yale Center for Research Computing

Agenda

- **The Problem:** why pipelines?
- **Pipelining concepts**
- **An example workflow**
- **Bash & Slurm**
- **Snakemake:** turning our example into a Snakemake pipeline
- **Break:** 10minute break
- **Nextflow:** using pipelines from the research community
- **Resources**

Setup

Log in to the cluster and clone the workshop repository:

```
git clone https://github.com/ycrc/pipelines-workshop.git
cd pipelines-workshop
ls examples/
```

You will need a terminal and a text editor.

We recommend an [Open OnDemand](#) VS Code session.

The Problem

Your Workflow

- Multiple steps that process input to produce output
- Some steps depend on others completing first
- It works: now you need to run it many times, scale it up, or share it

```
# step 1: process raw data
./clean.sh raw.dat > clean.dat

# step 2: run analysis
./analyze.sh clean.dat > results.dat

# step 3: make figures
./plot.sh results.dat > fig.png
```

What Can Go Wrong

- Script versions multiply
- Data folders accumulate
- "It worked on my machine"
- A step fails halfway: is the output valid?


Today's Learning Goals

- Understand key concepts for constructing data pipelines
- Build a simple workflow using bash scripts and Slurm
- Translate that workflow into a Snakemake pipeline
- Run a community-maintained pipeline using Nextflow and nf-core

Pipelining Concepts

Flowcharts and DAGs

- A workflow is a **directed acyclic graph** (DAG)
- Nodes are tasks, edges are dependencies
- No cycles: a task can't depend on its own output



Atomicity

- Every step of a pipeline should be **atomic**: it either fully succeeds, or fully fails.
- If a step fails, it should not produce partial output
- Prevents downstream steps from running on bad data

Reproducibility

- **Same input + same options = same output**
- Portable: works the same on any system
- Version control your pipeline, not just your analysis
- Pipelining tools have features to log exactly what processing was run in what order, with what parameters.

Our Example Workflow

The Input Data

- 10 plays by William Shakespeare
- UTF-8 plaintext files
- Stand-in for your real data: genomic reads, simulation output, etc.
- Small enough to run in a workshop, but the tools scale

The Goal


- Compute a measure of similarity between each pair of plays based on their most common words.

The Workflow

1. **Clean** each play (lowercase, remove punctuation)
2. **Count** word frequencies
3. **Extract** top 100 words per play
4. **Compare** every pair of plays (Jaccard similarity)
5. **Combine** into a similarity matrix CSV


The DAG (Simplified)

For two plays, the workflow looks like this:



The DAG (Full)

With all 10 plays, the DAG fans out: 45 compare steps:



The Bash Scripts


Our original scripts are found in the workshop repository under `examples/bash/` :

Script	Purpose
<code>01_analyze_play.sh</code>	Clean text, count words, extract top 100
<code>02_compare_plays.sh</code>	Jaccard similarity between two plays
<code>03_combine_results.sh</code>	Aggregate results into CSV
<code>00_run_all.sh</code>	Run everything in order

01_analyze_play.sh

Takes one play, produces its top 100 most frequent words:


clean text → count words → extract top 100



02_compare_plays.sh

Takes two plays, computes **Jaccard similarity** of their top-100 word sets:

Jaccard = $|\text{intersection}| / |\text{union}|$



03_combine_results.sh

Loops through all `.similarity` files, builds a final CSV: `output/similarity_matrix.csv`

00_run_all.sh: The Orchestrator

```
# Step 1: Analyze all plays
for play in data/*.txt; do
    name=$(basename "$play" .txt)
    ./analyze_play.sh "$name"
done

# Step 2: Compare all pairs
plays=(data/*.txt)
for ((i=0; i<${#plays[@]}; i++)); do
    for ((j=i+1; j<${#plays[@]}; j++)); do
        ./compare_plays.sh "$name1" "$name2"
    done
done

# Step 3: Combine results
./combine_results.sh
```

What's Wrong With This?

- Runs everything **serially**: no parallelism
- No **dependency tracking**: if one step fails, downstream runs anyway
- No **checkpointing**: must restart from scratch on failure
- **Manual cleanup** of intermediate files

Hands-On: Run the Bash Pipeline

1. `cd examples/bash`
2. `bash run_pipeline.sh`
3. Check `output/similarity_matrix.csv`
4. Simulate a data change: `echo "change" >> ../data/hamlet.txt`
5. Re-run: `bash run_pipeline.sh`
6. Notice the **entire** pipeline runs again: even the 9 unchanged plays

This naïve approach to our pipeline has obvious drawbacks...

Moving to Slurm

Our script works, but we're running it on the login node.
For a real workflow, we commonly:

- **Request dedicated resources:** CPU, memory, time
- **Run in the background:** submit the job and come back later
- **Get notified:** email when the job finishes or fails

We can wrap `00_run_all.sh` in a Slurm job script with `#SBATCH` directives. This is better, but still a single serial job: no parallelism.

Snakemake

What is Snakemake?

- Python-based workflow management tool
- Define **rules** with inputs, outputs, and commands that produce output from input.
- Snakemake builds the DAG and runs tasks in the right order.
- Snakemake allows you to run shell code, or Python code in your scripts.

Key Concepts

- **Snakefile:** The main file that defines the workflow
- **Rules:** Define a single step in the pipeline
 - Has an `input` , `output` , and a `shell` element.
- **Wildcards:** Create input lists from filename patterns

The Default Target: `rule all`

Snakemake works **backwards** from a target. `rule all` declares what the pipeline should produce:

```
rule all:
    input:
        "output/similarity_matrix.csv"
```

- This is always the **first rule** in the Snakefile
- Snakemake traces dependencies backwards to figure out what needs to run
- Nothing runs unless it's needed to produce this target

Translating: Clean Text

Bash

```
cat "$INPUT" \
| tr '[:upper:]' '[:lower:]' \
| tr -d '[:punct:]' \
| tr -s '[:space:]' '\n' \
> output/${PLAY}.clean.txt
```

Snakemake

```
rule clean_text:
    input:
        "data/{play}.txt"
    output:
        temp("output/{play}.clean.txt")
    shell:
        """
        cat {input} \
        | tr '[:upper:]' '[:lower:]' \
        | tr -d '[:punct:]' \
        | tr -s '[:space:]' '\n' \
        > {output}
        """
```

- `{play}` is a **wildcard**: one rule handles all 10 plays

Y|CRC `temp()` marks the file for automatic cleanup

Translating: Count Words

Bash

```
cat output/${PLAY}.clean.txt \  
| sort \  
| uniq -c \  
| sort -rn \  
> output/${PLAY}.counts.txt
```

Snakemake

```
rule count_words:  
    input:  
        "output/{play}.clean.txt"  
    output:  
        temp("output/{play}.counts.txt")  
    shell:  
        """  
        sort {input} \  
        | uniq -c \  
        | sort -rn > {output}  
        """
```

- Snakemake knows `count_words` depends on `clean_text` because the **output** of one matches the input of the other

Translating: Top 100 Words

Bash

```
head -100 \
  output/${PLAY}.counts.txt \
  > output/${PLAY}.top100.txt

rm output/${PLAY}.clean.txt
rm output/${PLAY}.counts.txt
```

Snakemake

```
rule top_words:
    input:
        "output/{play}.counts.txt"
    output:
        "output/{play}.top100.txt"
    shell:
        """
        head -100 {input} > {output}
        """
```

- No manual `rm` needed: `temp()` files are deleted after all rules that need them have finished
- `top100.txt` is **not** `temp()` because it's a final deliverable we want to keep

Translating: Compare Plays

Bash

```
comm -12 \
  <(awk '{print $2}' "$FILE1" \
    | sort) \
  <(awk '{print $2}' "$FILE2" \
    | sort) \
  > output/common.txt
# ... compute Jaccard ...
```

Snakemake

```
rule compare_plays:
    input:
        top1="output/{play1}.top100.txt",
        top2="output/{play2}.top100.txt"
    output:
        "output/{play1}_{play2}.similarity"
    shell:
        """
        COMMON=$(comm -12 \
          <(awk '{{print $2}}' \
            {input.top1} | sort) \
          <(awk '{{print $2}}' \
            {input.top2} | sort) \
          | wc -l)
        """
```

Y|CRC Two wildcards `{play1}` and `{play2}` handle all 45 pairs

May 26, 2026

Translating: Combine Results

This rule needs to know about **all** pair combinations upfront. We build the list at the top of the Snakefile:

```
# At the top of the Snakefile:
(PLAYS,) = glob_wildcards("data/{play}.txt")
PAIRS = []
for i, p1 in enumerate(PLAYS):
    for p2 in PLAYS[i + 1 :]:
        PAIRS.append((p1, p2))
```

- The loop generates all 45 pairs of input files automatically

Translating: Combine Results

```
rule combine_results:
    input:
        [f"output/{p1}_{p2}.similarity"
         for p1, p2 in PAIRS]
    output:
        "output/similarity_matrix.csv"
    shell:
        """
        echo "play1,play2,similarity" > {output}
        for file in {input}; do
            # parse filename, append row
        done
        """
```

Running Snakemake

When executing `snakemake`, it will find a `Snakefile` in the current directory.

- `snakemake -n` for a dry run
- `snakemake` to execute the pipeline
- `snakemake --dag | dot -Tpng > dag.png` to visualize


The Snakemake DAG

```
snakemake --dag | dot -Tpng > dag.png
```



The Snakemake DAG (Zoomed In)

A look at Macbeth/Othello:



What You Get "For Free"

- DAG visualization
- Automatic dependency resolution
- Only re-runs steps whose inputs changed
- Parallel (multiple processes) execution with `-j`
- Dry-run mode

Snakemake on Slurm

- `--executor slurm` : each rule becomes a separate Slurm job
- Snakemake monitors and schedules automatically

```
snakemake -j4 --executor slurm \
  --default-resources slurm_partition=day mem_mb=1000 cpus_per_task=1
```

Note: If running Snakemake from a Python environment, make sure you install `snakemake-executor-plugin-slurm`.

Keeping Small Steps Local

Not every rule needs its own Slurm job. Use `localrules` to run lightweight steps on the head node:

```
localrules: all, combine_results

rule all:
    input: "output/similarity_matrix.csv"

rule combine_results:
    ...
```

- Everything else still gets submitted to Slurm via `--executor slurm`

Custom Resources Per Rule

Different steps have different needs. Use `resources:` to request more for heavy rules:

```
rule star_align:
    input: "reads/{sample}.fastq"
    output: "aligned/{sample}.bam"
    threads: 8
    resources:
        mem_mb=32000,
        runtime=60 # minutes
    shell: "STAR --runThreadN {threads} ..."
```

- `threads:` sets `--cpus-per-task` in the Slurm job
- `resources:` sets memory, runtime, and other Slurm directives
- `--default-resources` on the command line applies to rules that don't specify

Python Scripts

Use `script:` to call a Python script. Snakemake passes inputs/outputs automatically:

Snakefile

```
rule plot_results:
    input:
        "output/similarity_matrix.csv"
    output:
        "output/heatmap.png"
    script:
        "scripts/plot_heatmap.py"
```

scripts/plot_heatmap.py

```
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv(snakemake.input[0])
# ... build heatmap ...
plt.savefig(snakemake.output[0])
```

R Scripts

Same pattern for R: Snakemake provides a `snakemake@` object:

Snakefile

```
rule fit_model:
    input:
        "output/counts.csv"
    output:
        "output/model.rds"
    script:
        "scripts/fit_model.R"
```

scripts/fit_model.R

```
counts <- read.csv(snakemake@input[[1]])
# ... fit model ...
saveRDS(model, snakemake@output[[1]])
```

`shell:` still works for any command-line tool: MATLAB, Julia, etc.

Containers in Snakemake

Your pipeline needs specific software versions: containers bundle everything so results are reproducible anywhere.

Snakemake supports per-rule `container:` directives:

```
rule align_reads:
    input: "reads.fastq"
    output: "aligned.bam"
    container: "docker://biocontainers/bwa:0.7.17"
    shell: "bwa mem {input} > {output}"
```

You can also point to a local `.sif` file: `container: "/path/to/bwa.sif"`

On our clusters we use **Apptainer** (not Docker), but Snakemake will convert.

Running Snakemake with Containers

Enable with `--software-deployment-method apptainer` (or `--sdm`):

```
snakemake -j4 --sdm apptainer
```

- Snakemake auto-pulls and caches container images
- Each rule runs in its own isolated container
- Our Shakespeare example uses only basic shell tools so no container is needed, but real-world pipelines typically specify one per rule.

Hands-On: Snakemake

1. `cd examples/snakemake` and `module load snakemake`
2. Dry run: `snakemake -n`
3. Execute: `snakemake -j1`
4. Check: `head output/similarity_matrix.csv`
5. Simulate a data change and dry-run: only affected steps re-execute:

```
echo "change" >> ../data/hamlet.txt
snakemake -n    # 14 of 77 jobs will re-run
```

Demo: Snakemake on Slurm

A head job orchestrates, submitting each rule as a child Slurm job:

```
#!/bin/bash
#SBATCH --partition=devel
#SBATCH --time=00:10:00
#SBATCH --mem=1G
#SBATCH --output=pipeline.out

module load snakemake
snakemake -j2 --executor slurm --latency-wait 30 \
  --default-resources slurm_partition=day \
  mem_mb=1000 cpus_per_task=1 runtime=5
```

Break

10 minutes

Nextflow

What is Nextflow?

- Groovy-based workflow management
- **Processes** and **channels**
- Built-in container support (Docker, Apptainer)
- Dataflow programming model

Key Concepts

- **Processes:** define tasks with inputs, outputs, scripts
- **Channels:** connect processes, data flows through them
- **Operators:** transform and combine channels

Snakemake vs Nextflow

	Snakemake	Nextflow
Language	Python	Groovy
Approach	File-based (rules produce files)	Dataflow (channels pass data)
Learning curve	Lower (Python syntax)	Higher (Groovy + channels)
Config	Snakefile + config.yaml	nextflow.config + profiles
Community pipelines	Snakemake Catalog	nf-core

Snakemake is more intuitive, while Nextflow has more features for complex workflows.

Side by Side: clean_text

Snakemake

```
rule clean_text:
    input:
        "data/{play}.txt"
    output:
        temp("output/{play}.clean.txt")
    shell:
        """
        cat {input} \
        | tr '[:upper:]' '[:lower:]' \
        | tr -d '[:punct:]' \
        | tr -s '[:space:]' '\\n' \
        > {output}
        """
```

Nextflow

```
process clean_text {
    input:
        path play

    output:
        tuple val(play.baseName),
            path("${play.baseName}.clean.txt")

    script:
        """
        cat ${play} \
        | tr '[:upper:]' '[:lower:]' \
        | tr -d '[:punct:]' \
        | tr -s '[:space:]' '\\n' \
        > ${play.baseName}.clean.txt
        """
}
```

Nextflow: Channels and Workflow

Where does `path play` come from? The **workflow** block wires it up:

```
// Create a channel from all .txt files
plays_ch = Channel.fromPath("data/*.txt")

workflow {
    cleaned = clean_text(plays_ch)    // each file flows into clean_text
    counted = count_words(cleaned)    // output flows into count_words
    top100  = top_words(counted)      // and so on...
}
```

- A **channel** is a stream of data flowing between processes
- Nextflow automatically parallelizes: 10 files in the channel = 10 concurrent tasks

Hands-On: Nextflow Shakespeare

Our Shakespeare workflow is implemented in Nextflow in `examples/nextflow/shakespeare/`.

```
cd examples/nextflow/shakespeare
module load Nextflow
nextflow run main.nf -with-report shakespeare.html
```

When finished, inspect the output files in `output/` and the HTML report `shakespeare.html`.

Nextflow in Practice

Rather than focusing on our Shakespeare workflow, we'll look at the **most common real-world use case**: running an existing, community-maintained pipeline.

- Thousands of researchers use Nextflow this way every day
- Someone has already written, tested, and optimized the pipeline
- You provide your data and configuration: Nextflow does the rest

Nextflow Configuration

Configuration is separate from the pipeline code, in `nextflow.config`:

- **Profiles:** switch between environments (local, Slurm)
- On our cluster, we use the `apptainer` profile for containers

```
// nextflow.config example for Slurm
process {
    executor = 'slurm'
    queue    = 'day'
}
apptainer {
    enabled = true
    cacheDir = '~/scratch/apptainer_cache'
}
```

Containers in Nextflow

Like Snakemake, Nextflow supports per-process `container` directives:

```
process align_reads {
    container 'docker://biocontainers/bwa:0.7.17'

    input:
    path reads

    output:
    path "aligned.bam"

    script:
    "bwa mem ${reads} > aligned.bam"
}
```

You can also point to a local `.sif` file: `container '/path/to/bwa.sif'`

Running Nextflow with Apptainer

Enable in `nextflow.config` or use a profile:

```
nextflow run main.nf -profile apptainer
```

- Nextflow auto-pulls Docker images and converts them to Apptainer format
- Set `cacheDir` to scratch storage — images can be large
- This is already configured when you use `-profile apptainer` with nf-core pipelines

The `work/` Directory

Nextflow stores all intermediate files in `work/`:

```
work/
├── 86/9b3800...    # clean_text (hamlet)
├── 94/d89dc0...    # count_words (hamlet)
├── 6d/5c8b4b...    # top_words (hamlet)
└── ...
```

- Each task gets a unique subdirectory (hash-based)
- Inputs are **symlinked** in, outputs are written here
- `publishDir` copies final results out of `work/` to where you want them
- `work/` can get **very large**: use `nextflow run ... -w ~/scratch/` to put it on scratch storage.

- resume: Pick Up Where You Left Off

```
nextflow run main.nf -resume
```

- Nextflow checks `work/`: if a task's inputs haven't changed, it **skips** it
- Critical for long pipelines: a failure at step 180 of 194 doesn't mean starting over
- Also useful during development: change one process, re-run, only that step re-executes

Cleanup when you're done:


```
nextflow clean -f          # delete all work/ contents
```

NF-Core Pipelines

What is nf-core?

- Community of **100+ curated Nextflow pipelines**
- Standardized structure: every pipeline works the same way
- Containerized: all software dependencies bundled
- Tested and documented by active maintainers
- Browse pipelines at <https://nf-co.re/pipelines>

What is nf-core?



The screenshot shows the nf-core Pipelines website. At the top, there is a dark header with the nf-core logo, a search bar, and a menu icon. Below the header is a green banner with the word "Pipelines" in large white text, followed by the text "Browse the 144 pipelines that are currently available as part of nf-core." Below the banner is a dark section with a search bar, a "Status" dropdown, a "Stars" dropdown, and two icons (a grid and a list). Below this is a table of pipelines.

↕ Name	Description	Released	↕ Stars	↕ Last release
rnaseq	RNA sequencing analysis pipeline using STAR, RSEM, HISAT2 or Salmon with gene/ isoform counts and extensive quality control.	✓	1213	3.23.0
sarek	Analysis pipeline to detect germline or somatic variants (pre-processing, variant calling and annotation) from WGS / targeted sequencing	✓	552	3.8.1
scrnaseq	Single-cell RNA-Seq pipeline for barcode-based protocols such as 10x, DropSeq or SmartSeq, offering a variety of aligners and empty-droplet detection	✓	313	4.1.0

What is nf-core?

nf-core/rnaseq

RNA sequencing analysis pipeline using STAR, RSEM, HISAT2 or Salmon with gene/isoform counts and extensive quality control.

[Launch version 3.23.0](#)

<https://github.com/nf-core/rnaseq>

Introduction Usage Parameters Output Results Releases 3.23.0

Introduction

nf-core/rnaseq is a bioinformatics pipeline that can be used to analyse RNA sequencing data obtained from organisms with a reference genome and annotation. It takes a samplesheet with FASTQ files or pre-aligned BAM files as input, performs quality control (QC), trimming and (pseudo-)alignment, and produces a gene expression matrix and extensive QC report.

Workflow Diagram:

- Pre-processing:** FastQC, MultiQC, FastQC, MultiQC, FastQC, MultiQC, FastQC, MultiQC.
- Genome alignment & quantification:** STAR, HISAT2, Salmon, RSEM, Salmon, RSEM, Salmon, RSEM.
- Pseudo-alignment & quantification:** Salmon, RSEM, Salmon, RSEM.
- Quality control & reporting:** MultiQC, RSEM, Salmon, RSEM, Salmon, RSEM.

run with

- nextflow run nf-core/rnaseq
- nf-core
- Nextflow
- Seera Platform

video introduction

nf-core/rnaseq

February 6, 2022 1 pm CET

bytesize

Bit-size talks, gigabyte-sized science

subscribers 163 stars 1213

Why Use Pre-Built Pipelines?

- **Tested by hundreds of users:** all open-source!
- **Reproducible out of the box:** containers, pinned versions
- **Saves development time:** focus on your science
- **Consistent interface:**

```
nextflow run nf-core/<pipeline> -profile test,apptainer --outdir results
```

nf-core/rnaseq

The most widely-used nf-core pipeline: bulk RNA-seq analysis.

Steps:

1. **FastQC**: raw read quality check
2. **Trim Galore**: adapter and quality trimming
3. **STAR**: align reads to reference genome
4. **Salmon**: quantify gene expression
5. **MultiQC**: aggregate QC into one report

Test profile uses a tiny yeast dataset (~50K reads).

Nextflow on Slurm: Head Job Pattern

Submit a lightweight batch job that acts as the Nextflow head node. It orchestrates the pipeline and submits each task as a separate Slurm job:

```
#!/bin/bash
#SBATCH --partition=day
#SBATCH --time=02:00:00
#SBATCH --cpus-per-task=2
#SBATCH --mem-per-cpu=4G

module load Nextflow/24.10.2
nextflow run nf-core/rnaseq -r 3.14.0 \
    -profile test,apptainer \
    -c nextflow.config --outdir output
```

`sbatch run_rnaseq_slurm.sh` — head job submits child Slurm jobs for each process.

Mixing Local and Slurm Execution

Use `nextflow.config` to route heavy steps to Slurm and keep lightweight steps local:

```
process {
    executor = 'slurm'
    queue = 'day'
    time = '1h'

    // Keep lightweight steps local to avoid queue wait times
    withName: 'MULTIQC|CUSTOM_DUMPSOFTWAREVERSIONS' {
        executor = 'local'
    }
}
```

- `withName:` selects processes by name — use `|` to match multiple
- Small steps run on the head node instantly instead of waiting in the queue

Demo: Running nf-core/rnaseq

```
salloc --mem-per-cpu=4G --cpus-per-task=4
nextflow run nf-core/rnaseq -r 3.14.0 -profile test,apptainer --outdir output
```

Flag	Purpose
nf-core/rnaseq	Pull and run the pipeline from nf-core
-r 3.14.0	Pin to a specific pipeline version
-profile test	Use built-in test dataset (yeast)
-profile apptainer	Use Apptainer containers
--outdir output	Where to write output

Inspecting Output

```
output/
├── multiqc/           # HTML summary report
├── star_salmon/       # Aligned reads + quantification
├── fastqc/            # Per-sample QC reports
├── trimalore/         # Trimmed reads
└── pipeline_info/    # Execution timeline, versions
```

Open `output/multiqc/multiqc_report.html` for alignment rates, read quality, and gene detection at a glance.

Finding Pipelines for Your Research

Browse <https://nf-co.re/pipelines>, examples:

Domain	Pipeline
RNA-seq	nf-core/rnaseq
Variant calling	nf-core/sarek
Single-cell	nf-core/scrnaseq
ATAC-seq	nf-core/atacseq
Amplicon (16S)	nf-core/ampliseq
Fetch public data	nf-core/fetchngs

Finding Pipelines for Your Research

Snakemake has a pipeline catalog too, <https://snakemake.github.io/snakemake-workflow-catalog>:

Domain	Pipeline
RNA-seq	rna-seq-star-deseq2
Variant calling	dna-seq-gatk-variant-calling
Single-cell	epigen/scrnaseq_processing_seurat
ATAC-seq	epigen/atacseq_pipeline
Amplicon (16S)	cmc-aau/nanopore_16Samp
Fetch public data	epigen/fetch_ngs

Resources & Next Steps

- [nf-core documentation](#)
- [Nextflow training](#)
- [Snakemake documentation](#)
- [Yale HPC documentation](#) and office hours

Workshop Feedback

Please help us improve this workshop by sharing feedback via a 2-minute anonymous survey. Thank you.

For access: click the link or scan the QR Code:

https://yalesurvey.ca1.qualtrics.com/jfe/form/SV_ac86jTriewu9l8W



Questions?

Thank you!

Any remaining time can be used for additional questions and office hours.